

SHIV RAJ BANGAD

513-837-5613 | sbangad@purdue.edu | linkedin.com/in/shivrajbangad | srbngd.github.io/portfolio/

PROFILE

Business Intelligence Engineer with experience building scalable data pipelines, ETL workflows, and analytics systems using Python, SQL, PySpark, and cloud platforms. Experienced in translating large and complex datasets into reliable, analytics-ready data products and actionable business insights through automation and data-driven solutions.

EDUCATION

Purdue University, Daniels School of Business **West Lafayette, IN**
Master of Science in Business Analytics and Information Management **Dec 2026**

- **Relevant Courses:** Business Analytics, Data Mining, Machine Learning, Reinforcement Learning, Cloud Computing, Advanced Database Management, Spreadsheet Modelling, Computational Finance
- **Certifications:** AWS Certified Cloud Practitioner, AI-900 Microsoft Azure AI Fundamentals

University of Cincinnati, Carl H. Lindner School of Business (Double Major) **Cincinnati, OH**
Bachelor of Science in Business Analytics and Bachelor of Business Administration in Finance **May 2025**

- **GPA:** 3.9/4.0
- **Achievements:** *Lindner Business Honors*: Kolodzik Business Scholar (Awarded to the top 5% of business students)

PROFESSIONAL EXPERIENCE

Amazon (AWS Analytics Engineer) **Seattle, WA**
Business Intelligence Engineer Intern **May 2026 - Aug 2026**

- Engineered an AI agent system to automate feature adoption tracking, cutting dashboard delivery from 2 days to 30 minutes
- Architected an automated PySpark/S3 ETL pipeline supporting 10B+ rows across 100+ AWS service features
- Programmed deterministic AST SQL guardrails on AWS Lambda to validate 100% of LLM, eliminating hallucinated joins
- Deployed an NL-to-SQL Q&A engine with automated anomaly detection to generate weekly performance narratives

Google (Youtube) **West Lafayette, IN**
Data Science Consultant (Capstone Project) **Jan 2026 - May 2026**

- Architected a recommendation engine using YouTube metadata to optimize content engagement for mid-tier creators
- Engineered a scalable framework to identify cross-channel semantic themes across 500+ channels to drive video ideation

BDT & MSD Partners (Private Equity) **Chicago, IL**
Financial Analyst Intern **Jun 2025 - Aug 2025**

- Prepared and reconciled ~\$100M in capital calls and distributions, ensuring 100% accuracy in investor allocations
- Automated cash workflows with SQL queries and Excel models, cutting processing time by 30%

University of Cincinnati - Center for Business Analytics **Cincinnati, OH**
Data Analyst (Cintas) **Jan 2025 - May 2025**

- Architected an automated anomaly detection dashboard using Python, cutting manual audit time by 360+ hours annually
- Leveraged feature engineering and statistical normalization to boost model predictive performance by 30%

Data Analyst (HBH Holdings) **Sep 2024 - Dec 2024**

- Built a Python pipeline to convert legal PDFs/Word files into structured text using Poppler, and NFKC normalization
- Delivered a Haystack and Mistral 7B contract Q&A assistant, boosting search accuracy by 75% with sub-5 second retrieval

Great American Insurance Group **Cincinnati, OH**
Data Analyst Intern **Apr 2022 - Apr 2024**

- Executed monthly data processing for 70,000+ policies, managing a total premium value of approximately \$1 billion
- Designed Access databases and SQL queries to streamline workflows, reducing retrieval times by 40%
- Developed an Excel-to-Tableau dashboard, automating reporting and reducing analysis time by 50%

PROJECTS

Data4Good National AI Competition **Oct 2025 - Jan 2026**

- Secured Regional Championship title by developing a custom DeBERTa model that achieved 97.5% accuracy
- Engineered a multi-class NLI system to classify semantic relationships as factual, contradictory, or irrelevant

Predictive Modelling for Firm Bankruptcy | Kaggle Competition **Nov 2025 - Dec 2025**

- Achieved 1st place out of 37 teams, securing a top-tier ROC-AUC of 0.9493 in predicting corporate insolvency risk
- Built a 3-stage super-learner combining XGBoost and LightGBM models with a Logistic Regression meta-learner
- Developed an "Evolutionary XGBoost" algorithm to generate and select high-interaction features over 100 generations

Stock Market Return Prediction **May 2024 - Aug 2024**

- Developed BART, SVM, and Random Forest models to predict stock market returns using macroeconomic/financial data
- Improved out-of-sample R^2 by 10% through advanced feature engineering and hyperparameter optimization

SKILLS

- **Languages & Cloud Data:** Python (PySpark, Pandas), Advanced SQL, R, AWS (S3, Glue, Athena, Lambda), Snowflake, PostgreSQL, MySQL
- **Machine Learning & AI:** Bedrock AgentCore, PyTorch, TensorFlow, LLMs, RAG Architectures, Vector DBs, NLP
- **BI & Stats:** Amazon QuickSight, Tableau, Power BI, ETL Automation, Dimensional Modeling, A/B Testing